Tactix: stable public Lucky actions across proof modes

Codex for Andrew

27 September 2026

Change and reason. The terminal Opus AM473 campaign found eight random public-program twins whose first difference was Lucky enumerating 19 actions with explicit DSR output and 21 without proof. There was no SAT/UNSAT contradiction, but the proof writer changed the sampled public action. Lucky now always enumerates stable IDs 0–20. Startup-PBP actions 19 and 20 still decline before mutation when the original-input boundary or selected proof format cannot certify them. Proof routing, public X7 legality and DefaultPolicy's explicit startup schedule are unchanged.

Focused validation. A public C-API regression compares proofless and strict-DSR enumeration, applies both unsupported startup actions and confirms an empty DSR stream. Release and proof-shadow Debug API, parity/count and random-policy tests pass. Both builds' parity and count startup PBPs independently verify against their original CNFs through VeriPB and CakePB; the missing-derivation count control is rejected. This is affected-scope evidence, not full N1/N2 acceptance.

Remaining work. Opus's 30 direct Proof API D2F flags (29 DSR, one DRAT) remain an open production-reach question. Those scripts can bypass kernel root-state maintenance; the public X7 validator checks live root units. Other producers emit transient proof lemmas, so their checker-visible lifecycle needs a targeted public reach audit before classifying D2F as out of contract. A new Opus campaign and full Debug run on this source are due, followed by the automatic-proof Rank/B28+C5, WV100 and WV2-89 gates. N1–N4 remain open; Tier D has not been submitted.

Evidence. `benchmarks/reports/2026-09-27-am482-stable-lucky-public-enumeration.md`; `benchmarks/reports/2026-09-27-am481-opus-terminal-debug-checkpoint.md`.