Tactix: current checker route and goal-exit checkpoint

Codex for Andrew

27 September 2026, 22:32 UTC

Current outcome. The user-approved native DSR contract is implemented and pushed on solver source `0a906338`: every native DSR first passes no-option `dsr-trim`, then a solver-declared final checker, CakeLPR for a witness-free elaboration and Trestle for a witnessed one. The production runner uses the solver declaration and scans proof content only to reject a mismatch. No legal X7 action or certificate requirement was narrowed. Full WV100 array `1731686_0-99` now has 99/100 final valid records: all 33 SAT models checked against the original CNF and 66 of 67 UNSAT proofs completely checked. The only open row is the current-source `ncc_none_21015` proof. Dependent scorer `1731826` waits for all 100 terminal tasks. This is an interim result, not Phase 2 acceptance or N1–N4 verification.

Evidence and resource contract. The array uses a frozen 100-row matrix, hashed inputs, one plain Release binary, seven checker pins, 12 CPUs and 32 GiB per task, a 1,200-second solver limit, and a 57,600-second checker window without overrides. No array throttle or node constraint was used. The long `snw_16_9` SAT tail passed in 973.254 seconds with a verified model. The 273 MB grid startup PBP passed VeriPB and CakePB. `miles1500.72` passed the full DSR-trim/LSR/Trestle route in 2,290.537 seconds, close to its older 2,299.631-second check. The current `ncc` proof is a different SHA from the archived 10,005-second DSR-trim/CakeLPR diagnostic; that old check does not certify this row. The 113 Python route/runtime/fuzz-protocol tests, focused Release/Debug proof serializer tests and current-source Debug `micro_proofs` passed. Current-source full Debug job `1731798` has 187/188 reported passes and the long scripted-default demonstration still running. The pre-route full Debug job passed 188/188 but is not substituted for the current-source job.

Goal-exit decisions. The user accepted complete Phase 1 Rank 50/50 and B28+C5 33/33 checked coverage with a 5% allowance on their aggregate solver-time targets: exact ceilings 8,521.52175 and 999.87615 seconds. The per-case 1,200-second/32-GiB limits remain. The older pre-route totals, 7,553.134 and 905.892 seconds, pass even the original stricter targets. The user also accepted WV100 100/100 and WV2–89 89/89 checked coverage at their 1,200-second and 600-second limits, respectively; the N1/N2 public reach and proof/model audit; and a separate online-learning readiness review. The Debug/fuzz/D2F acceptance item is parked while the user asks questions. The proposed per-tier Tier B/C parity and zero-open-N4-ledger gate is too strict for this deliverable; the user wants overall performance. A concrete full-union solved-count/PAR-2 replacement has been proposed for decision. No waived or unrun condition is called passed.

Next work and limitations. The held current-source WV2–89 bundle AM511 has all 89 input hashes, 17 pins, resource limits and scorer validated. It remains unsubmitted until AM507's terminal result is assessed. The last older-pin WV2–89 campaign covered only 78/89 within 600 seconds; an already-repaired checker failure accounts for one old miss, while the remaining rows include

SAT runtime gaps and winner-specific technique questions. The current-source full Debug run, the public D2F reach boundary, broad N1/N2 evidence, and the exact overall-performance threshold remain open. A prior learning architecture audit found opt-in inprocessing order/dose available, but production online learning not yet ready: a default-exact shared slot controller and owned safe live branching/deletion strategy handles are missing. A final readiness verdict belongs after the agreed gates, and can be "not ready" without invalidating the review.

After the preflight contract is settled and passed, the requested Tier D run must cover all 1,166 unique instances on `cpu-dense-hwulab-q`, at 5,000 solver seconds, four CPUs and 32 GiB per instance, with complete independent checking and cleanup. It has not been submitted. Full Tier D performance is a separate final measurement, not something that can be proved before that run.

Primary evidence. `benchmarks/reports/2026-09-27-am506-dsr-trim-first-gate.md`; `benchmarks/reports/2026-09-27-am507-wv100-dsr-trim-first-launched.md`; `benchmarks/reports/2026-09-27-am510-exit-gate-user-direction.md`; `benchmarks/reports/2026-09-27-am511-current-route-wv2-89-prepared.md`; `benchmarks/reports/2026-09-27-am512-wv100-dependent-score.md`; `benchmarks/reports/2026-09-25-am391-learning-readiness-current-source.md`.