Tactix: selected default gates recovered

Tactix Development Project

October 3, 2026

All four selected DefaultPolicy gates pass on unchanged source `375b621b`, using the single preregistered full WV100 recovery. The original failed WV100 campaign is preserved as rejected evidence. No source tuning, allowance increase or automatic retry loop was used. **Whole-goal acceptance remains pending full Debug and the owner's terminal both-policy fuzz/coverage report.**

Selected regression evidence

Set	Solver seconds	Checked	Checker TO	Solver TO
B28+C5	956.934	32	1	0
Rank 1–50	7,580.783	41	9	0
WV100 recovery	2,783.045	98	2	0
WV2–89	4,092.871	77	2	10 retained
Total	—	248	14	10

All 272 selected rows are present and pass the maintained scorer. Independent checks accept 123 SAT models and 125 UNSAT proofs. Fourteen clean checker timeouts remain censored, not verified; ten WV2 solver timeouts remain retained checkpoint exceptions. No selected result has a checker rejection.

Resources, input sets, binary/checker pins, solver caps and aggregate ceilings remain unchanged. Every complete proof pipeline retains 1,800 seconds; comparable baseline time t allows $t + \max(0.30t, 20\,\text{s})$. All 125 completed comparisons satisfy that bound. Thirteen censored comparisons remain within declared bounds; one already-censored baseline leaves its comparison inconclusive. No extended checking or reference-solver rerun was used.

Recovery array 1747089 has 98 COMPLETED/exit-zero tasks and two FAILED/exit-1 tasks from the allowed checker deadlines. Its watcher exits 1 while its maintained gate score passes. Previously lost task 95 is independently model-verified SAT in 1,031.247 seconds, with the same certificate hash as earlier complete solves. The original failed campaign remains rejected.

Correspondence and bounded diagnosis

All 262 completed selected baseline/candidate pairs have equal normalized printed work statistics and certificate hashes, also equal to foundation `12e3c4f0`. SAT model bytes are not compared by that analysis; candidate models are independently checked. This is scoped benchmark correspondence, not universal architectural equivalence.

Two opposite-order matched pairs on the lost instance give:

Pair / order	Predecessor seconds	Candidate seconds
1747087 / predecessor first	1,046.445	1,035.187
1747088 / candidate first	1,031.935	1,021.118

All four models verify independently; complete work and certificate hashes agree. Ratios are 0.989242 and 0.989518. Candidate repeat range is 1,021.118–1,035.187 seconds; predecessor range is 1,031.935–1,046.445. The reverse pair is COMPLETED/exit-zero, with both leg statuses zero and twenty pin bindings verified before and after. No large candidate work or instruction-cost increase is demonstrated. The original timeout's cause remains unproved because hardware counters and final search were incomplete. Two pairs do not establish statistical significance or general competitiveness. The registered follow-up allowance is consumed; no further repeat is queued.

The scheduler database was temporarily unavailable. Existing watchers retried the same IDs, then confirmed terminal accounting after recovery. No solver job was restarted. Earlier snapshots retain their as-recorded pending confirmation; the final terminal metadata closes that gap.

Source review and remaining acceptance

Current review binds all 28 operation groups and 49 original sites against 30 sources: 28 unchanged hashes, plus inspected feedback-description changes and Substitute owned-watch cleanup. Twelve prior source/contract and seventeen skill hashes still match. Locators do not establish runtime reach; exact R1–R13 owner minimums remain required.

Full Debug 1746924 remains live with 211 passing names and the long scripted test outstanding. Exact 212-test JUnit, terminal exits and 984 postflight bindings remain required. Fresh fuzz legs 1747082/1747083 are scheduler-complete with exit zero; targeted work and owner reporting remain active. No private campaign contents were inspected, and job exit does not prove fresh counts, effective reach, purity or replay acceptance.

Next: retain all 795 source bindings, finish exact Debug and delivered owner acceptance, then audit every original objective requirement and produce the final handoff. Scoped frontend overhead remains accepted in its own report; reward choices and learning remain deferred. No cleanup, unrelated cancellation or push occurred. See the technical report and hashed companion for exact scores, rows, source review and provenance.