Tactix winner-integration checkpoint
Certified affine rows, noL discovery, and regressions

Development report for Andrew
October 5, 2026

Outcome

The latest completed DefaultPolicy regression checkpoint is accepted on all four standing sets. B28+C5, Rank1--50, and WV100 retain every expected disposition; WV2--89 gains one checked solve, ramsey_4_4_19, relative to the preceding accepted checkpoint. There is no wrong verdict, checker rejection, or scorer failure. The complete Debug campaign passes 289/289.

This benchmark evidence is for immutable Release source 5eb58e52c and Debug source 64de46fe9. It predates the new certified affine and noL commits. The current development head adds a shared-controller path for exact GF(2) rows and hardens noL discovery, but has only focused and microproof evidence. Formal fixed-target coverage remains 1/44: the Ramsey row. No portfolio claim is made from the new code yet.

Campaign	Dispositions	Solver sum	Independent disposition
B28+C5	33/33	956.316 s	17 models, 15 proofs, one clean proof censor
Rank1--50	50/50	7,559.461 s	25 models, 17 proofs, eight clean proof censors
WV100	100/100	2,718.740 s	33 models, 65 proofs, two clean proof censors
WV2--89	89/89	3,990.643 s	48 models, 30 proofs, two clean censors, nine allowed timeouts
Debug	289/289	814.11 s	242 executables and 295 bindings reverified

All proof censors are clean expirations of the single 1,800-second regression checker allowance, rather than proof rejection. Completed comparisons remain within their row limits. Relative to accepted a9f1e14d4, solver sums improve by 2.739 s on B28+C5, 110.413 s on Rank, 43.504 s on WV100, and 97.171 s on WV2. WV2 solver timeouts fall from ten to nine because ramsey_4_4_19 now solves and checks UNSAT in 0.167 s.

Certified GF(2) integration

Commits 3787d6724, 2b2d7ecb8, and 405b56724 implement one vertical slice of the requested architecture. X9 AffineConsequence owns an immutable pre-publication clause view, reduces a bounded GF(2) basis, pages exact rows, and commits one selected row at a propagated root. Direct rows publish only missing parity CNF. Wide rows publish a compact conservative Boolean extension instead of invoking a private local-search solver.

Each wide transaction remaps authenticated source equations into a dense proof namespace, derives a balanced XOR network, and retains exactly the final definition clauses and signed root unit. Transition helpers absent from the final network remain in the monotone namespace but are fixed true by vacuous RAT units after an exact active-population check. Every retained definition records its output and operands, so reconstruction evaluates the XOR network topologically and omits the extensions.

The transaction records source occurrences, exact selected equations, row and RHS, proof additions/deletions, distinct RAT roles, introduced and proof-only helpers, projection, action identity, work, yield, and extent. Preparation is mutation-free. Source loss, stale ownership, unsupported proof format, interruption, or any proof/work/namespace/journal bound declines before publication. Proof-on and proof-off allocate the same namespace.

The same direct or extended payload is available through additive C API v9, deterministic replay, canonical DefaultPolicy execution, the opt-in gf2-affine random producer, and future learned-policy submission. The learner therefore chooses an ordinary typed action rather than entering a separate solver. M1 propagation and semantic-variable advice remain separate follow-ups.

Sixteen focused Debug tests pass. The exact two-source arity-eight fixture has 15 source variables, 256 input clauses, 13 retained definitions, and 17 retired helpers; its 972-lemma transaction proof is accepted by pinned drat-trim in SAT-forward mode. A contradictory width-eight certificate is accepted as a full refutation. Broader source overlaps, three-or-more-source rows, DSR checking of the wide path, real targets, full Debug, and the regressions remain pending for this new source.

Lucky/25 remains diagnostic. On noL-11-10 and noL-11-4, exact recognition fired repeatedly but both 600-second runs timed out after roughly 0.75 billion search moves, without a model or consumed partial guidance. This is a search and publication miss rather than a recognizer miss; it does not meet the Boolean integration objective.

Commit 8d13ca87a advances the certified route. ULC reauthenticates the checker-visible live population, requires the unique irredundant ALO source, chooses an unassigned endpoint, and emits pivot-first RAT through DSR and DRAT. X17 replaces combination enumeration with a canonical relation-clause index and durable bounded pages. Focused Debug and Release tests pass; a synthetic noL census finds 363 domains, 2,486 relations, 7,458 tuples, and 22,374 exact prefixes. The ULC microproof passes independent DSR and DRAT checking. The two real noL targets have not yet been recovered.

What the evidence permits

The accepted regression result says the pre-integration DefaultPolicy remains strong and the Ramsey Boolean transfer improved WV2 without a regression. The new affine/noL work establishes certified learner-selectable mechanics and focused proof soundness. It does not establish target effect or portfolio superiority. The fixed 36 and all 44 misses remain recovery obligations, with B28+C5, Rank1--50, WV100, WV2--89, full Debug, and both-policy fuzz as gates.

The next immutable checkpoint should run full Debug, then probe WV2-132 Tseitin, WV100-7 two-trees, WV100-51 MDP, rank 263 lightsout, and rank 569 lec_mult. Logs must distinguish absent structure, offer rejection, proof-work exhaustion, successful publication, and post-publication search effect. Only after coherent target evidence should the four full regressions be repeated.

The saved d170b4a5 comparison remains the approved authority and was not rerun. DSR remains the default proof format. A rule-specific pseudo-Boolean route remains an option only if measured compact clausal certification proves impractical.